

John Zheng  
jz5pt  
4/11/19  
inlab9

## Optimization

For the inlab, I chose to research the topic of optimization using the -O2 flag. I've included the sample code I wrote. The first function, helloWorld, just prints "Hello World!" to the console. The second function, positive, takes in an integer and returns true if the integer is positive. I call the functions in the main method. When I compile the code into assembly without the -O2 flag, the first thing I noticed was how long it was. It was 1582 lines just for the C++ code below. After I put the -O2 flag in, the assembly code was reduced to just 584 lines, still a lot but a big improvement. I've included the assembly code for the positive function below, the left one is the assembly with optimization. You can see the noticeable difference between the two. Without the optimization, the commands are really basic and there is a lot of mov commands. The optimized function uses the test command, which is essentially doing a bitwise AND and compare at the same time.

```
1 using namespace std;
2
3 #include <string>
4 #include <iostream>
5
6 void helloWorld(){
7     cout << "Hello World!" << endl;
8 }
9
10 bool positive(int x) {
11     return x > 0;
12 }
13
14 int main() {
15     helloWorld();
16     positive(7);
17
18     return 0;
19 }
```

```
118 .globl __Z8positivei
119 .p2align 4, 0x90
120 __Z8positivei:
121 .cfi_startproc
122 ## %bb.0:
123 push rbp
124 .cfi_def_cfa_offset 16
125 .cfi_offset rbp, -16
126 mov rbp, rsp
127 .cfi_def_cfa_register rbp
128 test edi, edi
129 setg al
130 pop rbp
131 ret
132 .cfi_endproc
133
```

```
76 .p2align 4, 0x90 ## -
basic_ostreamIT_T0_EES7_
77 __ZNSt3__14endlIcNS_11char_traitsIc
char_traitsIcEEEEENS_13basic_ostream
78 .cfi_startproc
79 ## %bb.0:
80 push rbp
81 .cfi_def_cfa_offset 16
82 .cfi_offset rbp, -16
83 mov rbp, rsp
84 .cfi_def_cfa_register rbp
85 sub rsp, 32
86 mov qword ptr [rbp - 8], rdi
87 mov rdi, qword ptr [rbp - 8]
88 mov rax, qword ptr [rbp - 8]
89 mov rcx, qword ptr [rax]
90 mov rcx, qword ptr [rcx - 24]
91 add rax, rcx
92 mov qword ptr [rbp - 16], rdi ##
93 mov rdi, rax
94 mov esi, 10
95 call __ZNKSt3__19basic_iosIcN
96 mov rdi, qword ptr [rbp - 16] ##
97 movsx esi, al
98 call __ZNSt3__113basic_ostrea
99 mov rdi, qword ptr [rbp - 8]
100 mov qword ptr [rbp - 24], rax ##
101 call __ZNSt3__113basic_ostrea
102 mov rcx, qword ptr [rbp - 8]
103 mov qword ptr [rbp - 32], rax ##
104 mov rax, rcx
105 add rsp, 32
106 pop rbp
107 ret
108 .cfi_endproc
109
```

Below is the code snippet for a product function that uses loops. Next to it are the code snippets in assembly, optimized in the middle. In this case, the compiler seems to know what the function's purpose is. Instead of using mov and conditional jumps, it just uses integer multiply on the first and the second parameter.

<pre> 14 int multiply(int x, int y) { 15     int sum = 0; 16     for (int i = 0; i &lt; y; i++) { 17         sum += x; 18     } 19     return sum; 20 } </pre>	<pre> 136 __Z8multiplyii: 137     .cfi_startproc 138     ## %bb.0: 139     push    rbp 140     .cfi_def_cfa_offset 16 141     .cfi_offset rbp, -16 142     mov rbp, rsp 143     .cfi_def_cfa_register rbp 144     mov dword ptr [rbp - 4], edi 145     mov dword ptr [rbp - 8], esi 146     mov dword ptr [rbp - 12], 0 147     mov dword ptr [rbp - 16], 0 148     LBB5_1: 149     mov eax, dword ptr [rbp - 16] 150     cmp eax, dword ptr [rbp - 8] 151     jge LBB5_4 152     mov eax, dword ptr [rbp - 4] 153     add eax, dword ptr [rbp - 12] 154     mov dword ptr [rbp - 12], eax 155     jmp LBB5_1 156 LBB5_4: 157     mov eax, dword ptr [rbp - 12] 158     pop rbp 159     ret 160     .cfi_endproc 161 </pre>
--	---

Overall, the main take away from this is that the non-optimized assembly has a lot of redundancy and the optimized assembly reduces redundancy by using more advance commands.